

1. (a) I decided that I wanted to keep the runtime more or less constant, so I used the following formula for Python:

$$SIZE = 2^i, \text{ and } ITERS = 2^{(25-i)}$$

- (b) I reran some of the points that were outliers, and chose the number that I got when I reran it because the rerun always ended up looking less like an outlier. (not very scientific though).
  - (c) See Figure 1 on the next page for Python data (I couldn't get the figures formatted correctly so I just put them on a separate page)
2. I changed my formula a little, because the C code was much much faster than the Python. I increased the number of iterations by  $2^5$ -fold:

$$SIZE = 2^i, \text{ and } ITERS = 2^{(30-i)}$$

3. After porting this to C and running it on different types, I found a few distinct patterns (see Figure 2 on page 3) The first thing that I noticed is the sharp increase in time per memory access after a given object size for each different datatype ( $2^{12}$  for int64 and double,  $2^{14}$  for int16, int32, and float, and  $2^{15}$  for int8, roughly). I concluded that this must be due to one of the memory caches filling up.

Another interesting trend is the way that these datatypes are grouped. I assume this is because of the amount of memory that each type takes up.

In addition, if you squint at the graph, you can almost see three levels, which I think may represent the 3 caches (but I could be wrong).

4. In order to determine the amount of memory used for each of the different C datatypes with  $2^{20}$  objects, we use the command `/usr/bin/time -v ./update 1048576 1`, and then read the "maximum resident set size (kbytes)" result.

datatype	memory usage (kb)
int64_t	50164
double	50128
float	25606
int32_t	25600
int16_t	13184
int8_t	7140




Figure 1: This is what Python looks like. It is slow

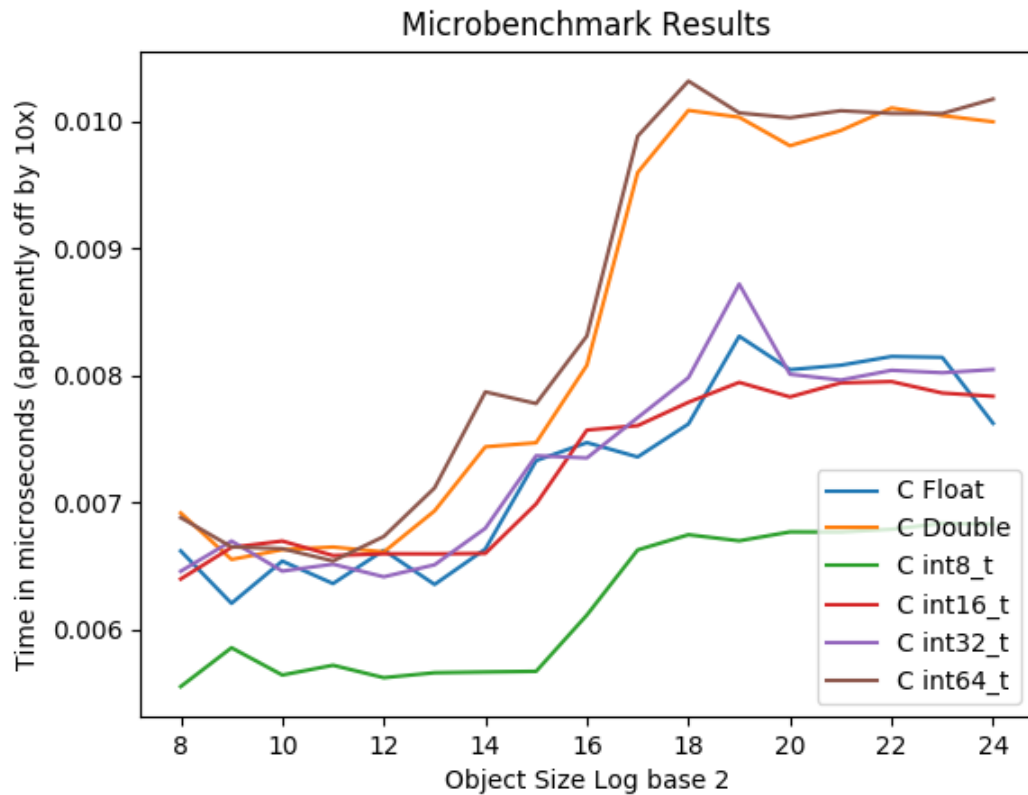


Figure 2: Look at the memory caches filling up. I wrote this on the y-axis label, but I will reiterate that apparently, according to people who know things about computer systems (Eitan and my mother) my measurements are 10x too fast to be feasible. I could not manage to find the bug in my calculations. I timed these with a stopwatch to see if I was going crazy and my timing calculations in the C code do not seem off by 10, according to the stopwatch.

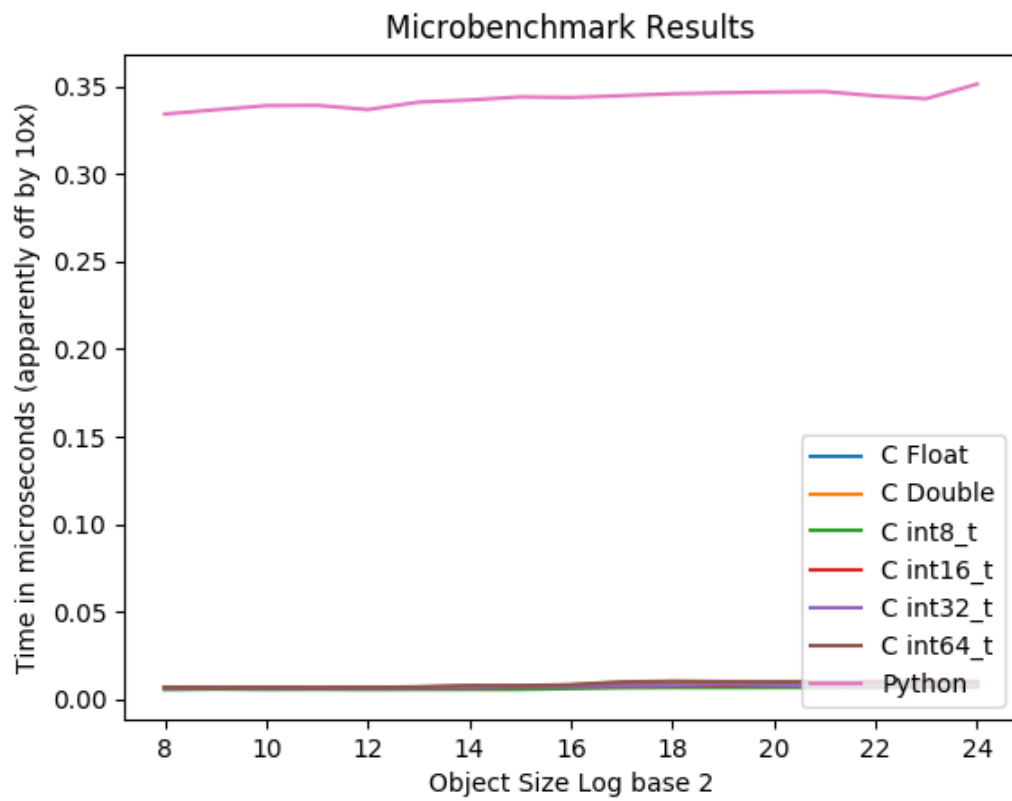


Figure 3: Look how slow Python is.